Intelligent Systems with Applications 16 (2022) 200141

Contents lists available at ScienceDirect

Intelligent Systems with Applications

journal homepage: www.journals.elsevier.com/intelligent-systems-with-applications

Automated MeSH term suggestion for effective query formulation in systematic reviews literature search☆

Shuai Wang [*,a], Harrisen Scells [a], Bevan Koopman [b], Guido Zuccon [a]

[a] The University of Queensland, Brisbane, Australia
[b] CSIRO, Brisbane, Australia

ARTICLE INFO

Keywords:
MeSH terms suggestion
Systematic review
Neural model
Evaluation

ABSTRACT

High-quality medical systematic reviews require comprehensive literature searches to ensure the recommendations and outcomes are sufficiently reliable. Indeed, searching for relevant medical literature is a key phase in constructing systematic reviews and often involves domain (medical researchers) and search (information specialists) experts in developing the search queries. Queries in this context are highly complex, based on Boolean logic, include free-text terms and index terms from standardised terminologies (e.g., the Medical Subject Headings (MeSH) thesaurus), and are difficult and time-consuming to build. The use of MeSH terms, in particular, has been shown to improve the quality of the search results. However, identifying the correct MeSH terms to include in a query is difficult: information experts are often unfamiliar with the MeSH database and unsure about the appropriateness of MeSH terms for a query. Naturally, the full value of the MeSH terminology is often not fully exploited.

This article investigates methods to suggest MeSH terms based on an initial Boolean query that includes only free-text terms. In this context, we devise lexical and pre-trained language models based methods. These methods promise to automatically identify highly effective MeSH terms for inclusion in a systematic review query. Our study contributes an empirical evaluation of several MeSH term suggestion methods. We further contribute an extensive analysis of MeSH term suggestions for each method and how these suggestions impact the effectiveness of Boolean queries.

1. Introduction

A medical systematic review is a comprehensive review of literature for a highly focused research question. Systematic reviews are seen as the highest form of evidence and are used extensively in healthcare decision making and clinical medical practice. In order to synthesise literature into a systematic review, a search must be undertaken. A major component of this search is a Boolean query. The Boolean query is often developed by a trained expert (i.e., an information specialist), who works closely with the research team to develop the search, and usually has some knowledge of the domain been searched. The most commonly used database for searching medical literature is PubMed. Due to the increasing size and scope of these databases, and of PubMed in particular, the Medical Subject Headings (MeSH) thesaurus was developed to conceptually index studies (Richter & Austin, 2012; Zieman & Bleich, 1997). MeSH is a controlled vocabulary thesaurus arranged in a hierarchical tree structure (specificity increases with depth in a parent→child relationship, e.g., `Anatomy→Body Regions→Head→Eye`... etc.). Indexing and categorising studies with MeSH terms enables queries to be developed which incorporate both free-text keywords *and* MeSH terms — enabling more effective searches. The use of MeSH terms in queries has been shown to be more effective than free-text keywords alone (Abdou & Savoy, 2008; Chang et al., 2006; Richter & Austin, 2012; Tenopir, 1985), e.g., they increase precision (Liu & Wacholder, 2017) and are far less ambiguous than free-text (Wacholder et al. 1997). However, it is still difficult even for expert information specialists to be familiar with the entire MeSH controlled vocabulary (Liu, 2009; Liu & Wacholder, 2017) — at the time of writing, MeSH contains 29,640 unique headings.

PubMed has attempted to overcome this difficulty by developing a

☆ This paper is currently in submission with Intelligent Systems with Applications Journal Technology-Assisted Review Systems Special issue and is under peer review.

* Corresponding author.
E-mail addresses: shuai.wang2@uq.edu.au (S. Wang), h.scells@uq.edu.au (H. Scells), bevan.koopman@csiro.au (B. Koopman), g.zuccon@uq.edu.au (G. Zuccon).

https://doi.org/10.1016/j.iswa.2022.200141
Received 8 June 2022; Received in revised form 29 August 2022; Accepted 11 October 2022
Available online 20 October 2022
2667-3053/© 2022 The Authors. Published by Elsevier Ltd. This is an open access article under the CC BY-NC-ND license (http://creativecommons.org/licenses/by-nc-nd/4.0/).

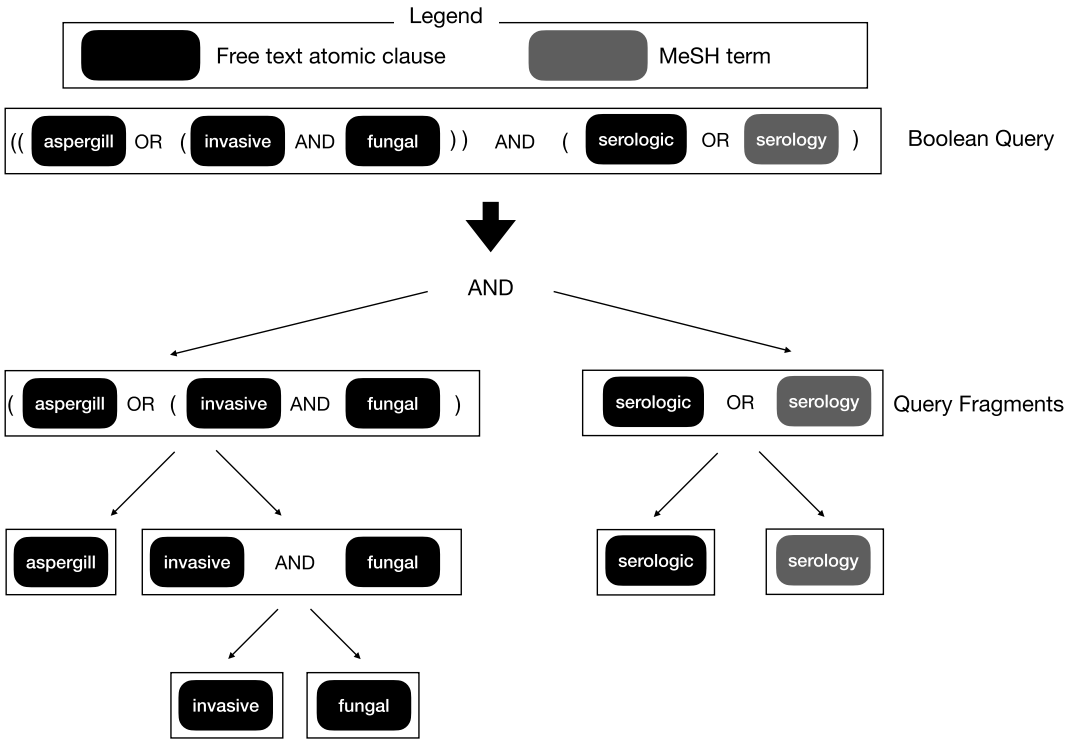

Fig. 1. Example query showing a **Boolean Query**, two **Query Fragments**, several **Free text atomic clauses**, and a **MeSH term**.

method called Automatic Term Mapping (ATM). ATM is an automatic query expansion method which attempts to seamlessly map free-text keywords in a query to one of the three categories (index tables): MeSH, journal name or author name (Nahin, 2003). Although ATM is applied by default for all queries issued to PubMed, it has several semantic limitations: it is inaccurate when used to expand free-text acronyms into MeSH terms (Schulz et al., 2001); it produces different MeSH expansions even though synonymic free-text terms are used (Adlassnig et al., 2009); and has difficulty disambiguating between MeSH terms and journal names (Smith, 2004). Despite these limitations, the use of ATM for MeSH term suggestion has been shown to increase the precision of free-text searches in the genomic domain (Lu et al., 2009), and is the state-of-the-art method for the MeSH Term suggestion task. However, its use has, to the best of the authors' knowledge, not been empirically evaluated in the context of improving the effectiveness of systematic review literature search queries.

Recent advances in the use of pre-trained language models (PLMs) such as BERT (Devlin et al., 2019), T5 (Raffel et al., 2019), and GPT-3 (Brown et al., 2020) have delivered state-of-the-art performance in many natural language processing tasks. Typically, a pre-trained language model is trained on a large corpus using the transformer architecture to "get familiar" with language representations. Then the model is fine-tuned to downstream tasks to perform with high effectiveness across the target task. The transformer architecture is an encoder-decoder model training structure that does not use recurrence and convolutions (Vaswani et al., 2017). Prior work showed that using PLMs can significantly increase effectiveness in ad-hoc search (Lin et al., 2021) as well as in professional search (Chalkidis et al., 2020; Choe et al., 2022; Qin et al., 2021; Yang et al., 2022).

In this article, we introduce the task of MeSH term suggestion for Boolean queries used in systematic review literature search[1]. We model this task within the context of an information specialist looking for MeSH terms to add to a query without MeSH terms currently present. We also propose a framework to evaluate the effectiveness of the suggestion of MeSH terms on established collections of systematic review literature search queries. This article adds to a recent stream of research that has focused on computational methods for the assisted formulation (Agosti et al., 2019; Scells et al., 2020a, 2021, 2020b) or refinement (Agosti et al., 2020; Alharbi & Stevenson, 2020; Harrisen & Guido, 2018; Scells et al., 2019; Wang et al., 2021a) of Boolean queries for systematic review creation, and more generally to research on computational methods for technology-assisted reviews (Cormack & Grossman, 2017; Lee & Sun, 2018, 2022; Li & Kanoulas, 2020; Sneyd & Stevenson, 2021). Furthermore, we propose two categories of methods for the MeSH term suggestion task, including methods based on the BERT pre-trained language model and methods not based on BERT (lexical methods). We show that our methods suggest MeSH terms that outperform the effectiveness of the MeSH terms selected by the information specialists and included in the original queries. Our methods are readily integrable into tools for information specialists to help with the construction of systematic review Boolean queries.

The contributions of this article are:

1. The introduction of the new task of suggesting MeSH terms for systematic review literature search (Boolean queries), modelled within the context of an information specialist looking for MeSH terms to add to a query without MeSH terms present.
2. The formulation of MeSH term suggestion methods to help information specialists and researchers to construct Boolean queries for systematic review creation.
3. An empirical evaluation of the effectiveness of different MeSH terms suggestion methods
4. An understanding of how the MeSH terms suggested by the proposed automatic methods differ from those originally selected by information specialists formulating the query.

[1] This article is an extension of our previous work published at the 2021 Australasian Document Computing Symposium Wang et al. (2021a).

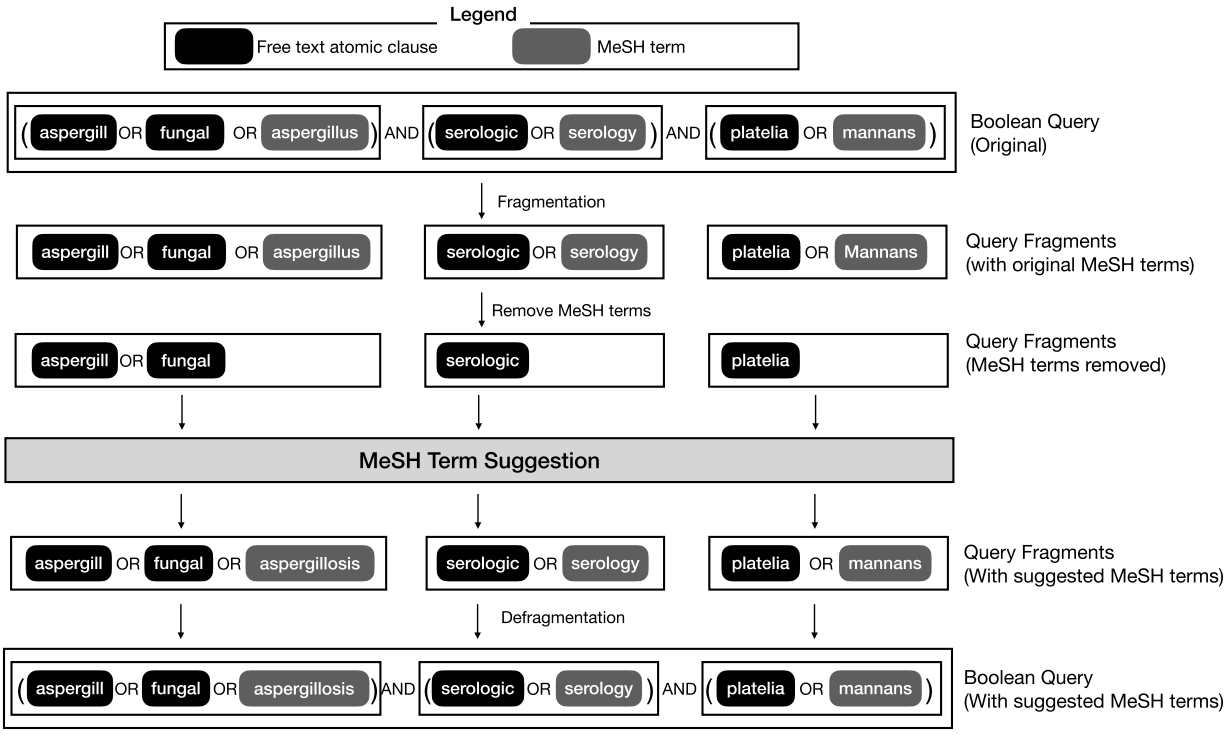

Fig. 2. Overview of the MeSH term suggestion procedure. Proposed methods using lexical MeSH term retrieval or BERT MeSH term retrieval facilitate the suggestion of MeSH terms. We evaluate each method that suggests MeSH terms in terms of (1) the ability for the suggested MeSH terms to effectively retrieve literature for a defragmented Boolean query, (2) overlap between suggested MeSH terms and MeSH terms included in the original query. Note that the number of MeSH terms suggested for a fragment may be lower or higher than the number of MeSH terms in the original query.

2. Material and methods

2.1. Overview of the MeSH term suggestion task

We start by outlining the task of MeSH term suggestion for Boolean queries that do not already contain MeSH terms. We assume the user has entered a Boolean query without MeSH terms. A Boolean query can be viewed as a tree where Boolean operators (e.g. AND, OR) represent the internal nodes of the tree, while free-text atomic clauses and MeSH Terms are the leaves. Free-text atomic clauses are one or more words that express a concept, e.g., a disease, a treatment or a population aspect. We call each of the first level nodes of the tree (i.e. the nodes at depth 1) a query fragment. Typically, a query fragment represents an individual aspect of an information need (Clark, 2013); specifically, each query fragment corresponds to a different PICO element, i.e. population, intervention, control, and outcome (Schardt et al., 2007). These concepts are shown in Fig. 1. The task of MeSH term suggestion is to identify appropriate MeSH terms to be added as leaves to a query fragment. In this article, we suggest MeSH terms for each query fragment independently of each other. We leave the investigation of query fragment dependencies concerning MeSH term suggestions for future work.

Fig. 2 gives an intuition for how we obtain query fragments from a Boolean query, how MeSH terms are suggested for a given query fragment, and how we perform *defragmentation* to construct a new Boolean query that includes MeSH terms. The figure shows that after fragmentation (i.e. the process of deriving query fragments), we remove all the MeSH terms from each query fragment. We then apply a MeSH term suggestion technique which adds new MeSH terms into a query fragment. The new query fragments that now contain suggested MeSH terms are then defragmented by combining all of the query fragments corresponding to the original query with the AND operator.

This work extends our existing line of research into MeSH term suggestion (Wang et al., 2021a), where we previously developed several techniques that depend on pre-existing lexical matching systems. One limitation of these systems is their dependence on manually crafted rules that are expensive to create and have limitations in terms of how words are matched to MeSH terms (e.g., spelling variants, acronyms, misspellings). This article instead investigates the use of pre-trained language models, i.e., BERT, for the task of MeSH term suggestion. These neural models have been shown to be resilient to the shortcomings of lexical-based systems (Devlin et al., 2019; Wang et al., 2021b). However, neural models have their own limitations, particularly requiring large amounts of training data. The following sections first provide a brief overview of our existing lexical-based techniques and then describe our new neural techniques in detail, specifically addressing the need for ad-hoc training data.

2.2. Lexical MeSH term suggestion

Our lexical-based methods are formulated as a pipeline of three steps: retrieval, ranking, and refinement. The following sections provide a brief overview of each of these steps. For a more comprehensive discussion of the lexical-based methods, refer to our previous work (Wang et al., 2021a).[2]

Retrieval

The first step in our MeSH term suggestionpipeline is the **retrieval** of MeSH terms. The retrieval of MeSH termsis facilitated by three different methods:

ATM The entire free-text only query fragment is submitted to the PubMed Entrez API (Sayers, 2010) for *automatic term mapping*

[2] Version 2018 with options set to default values.

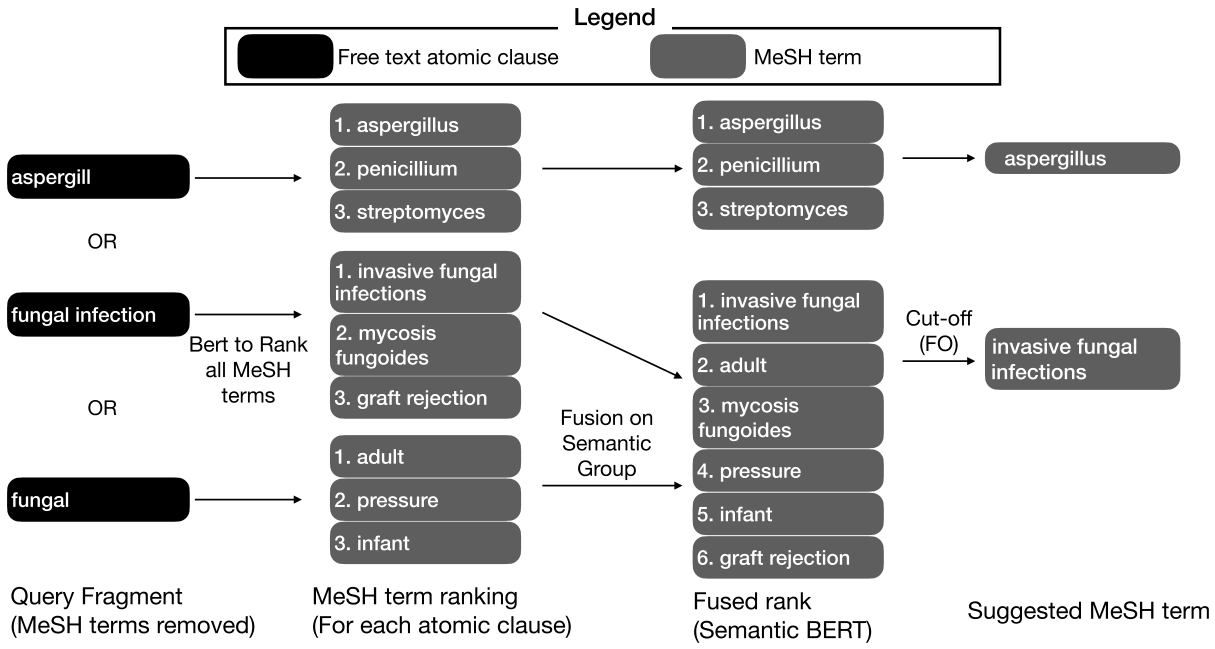

Fig. 3. Overview of the MeSH term suggestion for the BERT methods. Note that Fusion of MeSH ranks may be optional in the pipeline.

(ATM). This is the default system used by PubMed for automatically adding MeSH terms to queries.

MetaMap Each free-text atomic clause in a query fragment is submitted to MetaMap (Aronson, 2001).The results are filtered to only include those entities derived from the MeSH source. All of the mapped MeSH termsare recorded for each of the free-text terms in a query fragment. Additionally, the score is recorded for each MeSH term.

UMLS We index UMLS (Bodenreider, 2004) into Elasticsearch v7.6. Each free-text atomic clause in the query fragment with MeSH terms removed is submitted to the Elasticsearch index. The results are filtered to only include synonyms of concepts derived from the MeSH source. Additionally, the BM25 score is recorded for each MeSH term.

For the MetaMap and UMLS approaches, the same MeSH term may be retrieved multiple times for a given free-text fragment. To overcome this issue, we re-score the MeSH termsusing rank fusion (CombSUM) (Fox & Shaw, 1994). The intuition for this re-scoring is that highly common MeSH termsthat also obtain a high score from these retrieval methods should be scored highly overall (thus ranked higher than common MeSH terms *and* highly scoring MeSH terms).[3]

Ranking

Once MeSH terms have been retrieved, they are ranked according to the approach for entity ranking described by Jimmy et al. (2019) by adapting features proposed by Balog (2018). In total, we use eleven entity features. Positive instances correspond to MeSH termsin the original query fragment; negative instances correspond to MeSH termsnot in the original query fragment (binary labels). With features and instance labels, we train a learning-to-rank (LTR) model for each retrieval method. In addition to LTR, we also investigate a rank fusion approach (Fox & Shaw, 1994), where we combine the normalized MeSH term suggestion scores from each of the three methods to produce a new ranking that incorporates the highest ranking MeSH terms from each method. The intuition for investigating rank fusion in this context is that each method may retrieve different MeSH terms; and those terms may be ranked differently each time. Therefore, we boost MeSH terms that are retrieved and ranked highly by multiple methods.

Refinement Finally, we seek to refine the suggested MeSH terms by estimating a rank cut-off. We do this using a score-based gain function. Formally, the cumulative gain CG for a MeSH term at rank p is

$$CG_p = \sum_{i=1}^{p} score_i \tag{1}$$

where the score for a MeSH term is equal to $1 - normalised\ score$ (i.e., min-max normalisation) for the MeSH term.

We tune a parameter, κ, for each retrieval method which controls the percentage of total CG allowed to be observed before the ranking is cut-off (i.e., a refinement of the ranking). We tune κ from 5% to 95% in increments of 5%. The intuition for re-scoring MeSH terms becomes apparent when used with the κ parameter: the highest-ranking MeSH term will receive a score of 0, resulting in at least one MeSH term suggested for every query fragment.

Note that MeSH terms may share the same score, i.e., they may be tied. We take a conservative approach to account for the problem of tied MeSH terms at the boundary of the cut-off specified by κ. Whenever we encounter ties, we treat all of the tied MeSH terms as a single accumulation of gain that equals the summed gain across the scores of the tied MeSH terms. This treatment has the effect that tied MeSH terms account for much larger accumulations of gain. Therefore, tied MeSH terms at the top of rankings are more likely to be included in the cut-off than tied MeSH terms at the bottom. In essence, either all tied MeSH terms are considered within the cut-off (i.e., ties at the top of the ranking), or no tied MeSH terms are considered (i.e., ties at the bottom of the ranking).

2.3. BERT MeSH term suggestion

Next, we extend our MeSH term suggestion methods using fine-tuned PLM models. Firstly, PLM models are typically chosen from the same domain in which the task is conducted.

Architecture We show the architecture of our fine-tuning and inference processes in Fig. 4. We use BioBERT (Lee et al., 2020) as the base

[3] version 2019AB using the `MRCONSO`, `MRDEF`, `MRREL`, and `MRSTY` tables.

Figure 3: Overview of the MeSH term suggestion for the BERT methods. Note that Fusion of MeSH ranks may be optional in the pipeline.

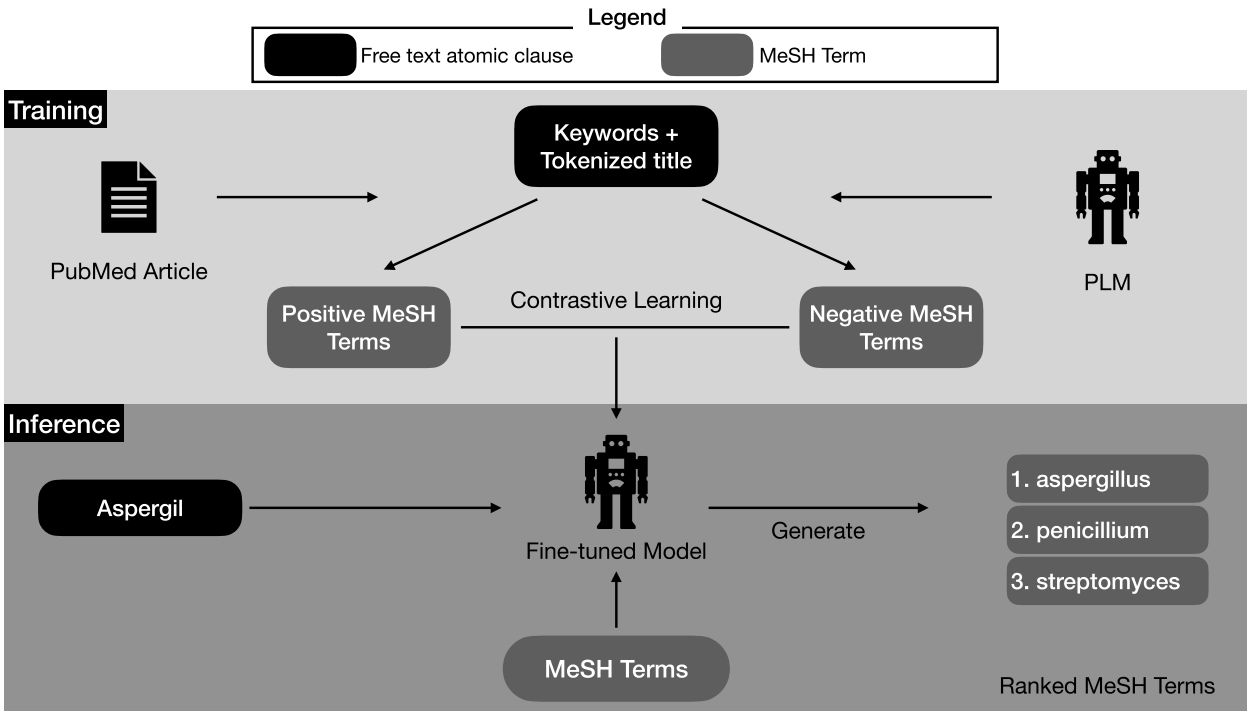

Fig. 4. Architecture of model fine-tuning and inference.

PLM, as the context of this paper is medical systematic reviews. BioBERT is a PLM pre-trained on PubMed abstracts and PubMed Central (PMC)[4] full-text articles using the BERTtraining architecture (Devlin et al., 2019). After fine-tuning, BioBERT has achieved state-of-the-art performance on many medical-related tasks, including biomedical named entity recognition, relation extraction and question answering (Lee et al., 2020).

Ideally, training data closely related to the target task should be used to fine-tune a PLM to achieve the highest effectiveness. Ideally, in our case, we would use professionally constructed medical systematic review Boolean queries to fine-tune our model. However, PLMs are typically data-hungry and require a large number of labelled training samples. In systematic review literature search, several public datasets are available with Boolean queries, such as the CLEF TAR collections (Kanoulas et al., 2017, 2019, 2018), the collection of Wang et al. (2022a), and the collection of Scells et al. (2017). Between these datasets, however, only 253 unique topics would be available to train the model: an insufficient amount to effectively fine-tune a BERT model.

Instead, we create training samples by approximating the target task using data obtained from PubMed. We use the publicly available PubMed baseline to obtain the metadata about all published articles up to the start of 2022. The metadata contains information such as the title and abstract, but importantly for this work, it also includes author-assigned keywords and the relevant MeSH terms for an article. We use the assigned keywords and MeSH terms for every article in the PubMed dataset to approximate the task of MeSH term suggestion. To maximise the amount of training data, we also extract keywords from the title (as not all PubMed articles contain keywords). To tokenise titles, we use the process described by Wang et al. (2022b). Firstly, we tokenise the title using Gensim (Khosrovian et al., 2008), and then we remove stopwords too using NLTK (Bird & Loper, 2004). We use the toolkit proposed by Gao et al. (2022) to develop a dense retriever to suggest MeSH Terms. The model is fine-tuned with localized contrastive loss using triples of $< k_{a,i}, m_a^+, m_a^- >$ where a is a PubMed article, $k_{a,i}$ is the ith keyword in the PubMed article, m_a^+ are the MeSH terms for the PubMed article, and m_a^- are ten randomly sampled MeSH terms from the MeSH thesaurus. Many MeSH terms contain spaces or punctuation. Our model considers each MeSH term a unique token in the model vocabulary. Once the model is fine-tuned, we obtain an encoding for all MeSH terms. At inference time, we create an encoding for a keyword to obtain a score using the `[CLS]` token for all MeSH terms. Thus, our method scores and ranks all MeSH terms given a keyword.

Ranking Suggestions The goal of MeSH term suggestion is to suggest MeSH terms for each query fragment. However, the result from the BERT suggestion method consists of a ranked list of MeSH term suggestionfor each free text atomic clause. We need to combine the rankings for each MeSH term. We formulate this combination task into two steps, (1) choosing how we represent a MeSH term ranking, and (2) choosing where to cut off the ranking. We present an overview of the combination task in Fig. 3.

First, we choose the best way to represent a ranking, which means deciding if MeSH terms should be suggested individually for every free text atomic clauses, as a whole for every fragment, or using other heuristics to decide how the representation should be computed. We designed three ranking representation methods:

1. **Atomic BERT**: Firstly, we treat suggestions for each free text atomic clauseindividually, essentially applying no strategy to combine the suggestions.
2. **Fragment BERT**: Next, we study the combination of all MeSH term rankings for a given query fragment. We apply rank fusion (normalised CombSUM Fox & Shaw, 1994) to all of the free text atomic clausesin a query fragment. For computational reasons, we only use the top 20 MeSH terms for each free text atomic clause.
3. **Semantic BERT**: Finally, we study semantically grouping free text atomic clausesand apply the same rank fusion technique as above,

[4] PubMed Central is the repository containing full-text articles of the open-access part of the PubMed database.

Table 1
Example query fragments with separation of semantic groups. In the example, 'neonatal sepsis', 'neonatal bacteremia' and 'neonatal infections' are grouped to form a semantic group, while 'death' is another semantic group.

MeSH Removed Fragment	neonatal sepsis OR neonatal bacteremia OR neonatal infections OR death			
free text atomic clauses	neonatal sepsis	neonatal bacteremia	neonatal infections	death
semantic group	neonatal sepsis, neonatal bacteremia, neonatal infections			death

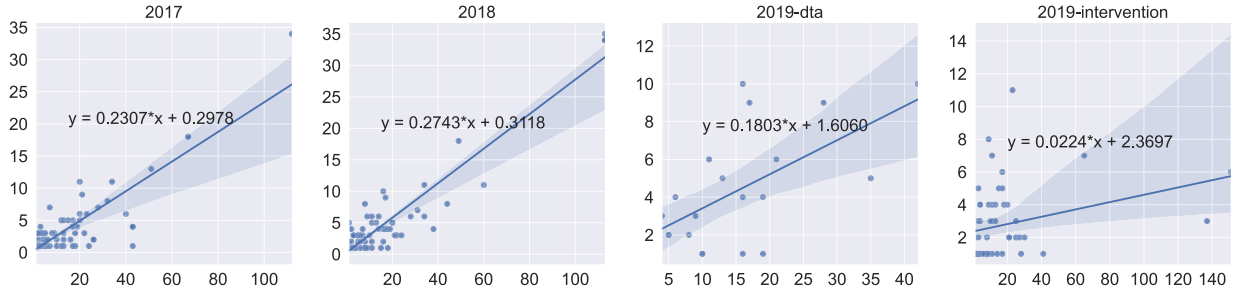

Fig. 5. Linear regression performed on the number of keywords (x-axis) and the number of MeSH terms (y-axis) in query fragments for training splits of CLEF TAR 2017, 2018, 2019-dta and 2019-intervention.

but this time to each group. We show an example of a semantic group in Table 1. To derive semantic groups, we first take all free text atomic clausesfrom the fragmentand obtain word2vec embeddings for each free text atomic clause. Then we compute cosine similarities between all free text atomic clauseto decide if they are semantically related. In our experiments, we apply a threshold of 0.7 on the similarity. We use a word2vec model pre-trained on PubMed and Wikipedia (Moen & Ananiadou, 2013). There are two reasons we use word2vec rather than BERT for semantic groups. First, if we apply our proposed BERT model, we note that we fine-tuned using semantic pairs of free text atomic clausesand MeSH terms: thus, calculating the similarity between two free text atomic clausescan cause a model mismatch. Secondly, the use of an additional BERT model will increase the latency in producing suggestions at inference time, as each free text atomic clauseneeds to be encoded twice.

Second, we choose where to cut off the ranking of MeSH terms from the ranking representations. We propose four strategies to cut off MeSH term rankings:

1. **First only (FO)**: The first MeSH term of the ranking is selected for each ranking representation.
2. **Same as free text atomic clauses(SA)**: The number of MeSH terms selected equals the number of free text atomic clausesin each fragment(i.e, only applicable to **Fragment BERT**).
3. **Same as original (SO)**: The MeSH terms selected equals the number of MeSH terms in the query fragment prior to removing MeSH terms (i.e, only applicable to **Fragment BERT**).
4. **Linear (LN)**: The number of MeSH terms selected is learnt using a linear function with respect to the number of free text atomic clausesin the fragments (i.e, only applicable to **Fragment BERT**).

2.4. Evaluation

The end goal of a systematic review literature search is to find all of the relevant literature at the minimum cost. Thus, an effective Boolean query minimises the number of documents retrieved while maximising the retrieval of relevant documents. In our MeSH term suggestion task, we use the retrieval effectiveness of defragmented Boolean queries to evaluate MeSH term suggestion.

The MeSH terms included in the original query have been derived often after careful consideration by expert information specialists. We therefore consider how the MeSH terms included in the original queries differ from those suggested by the methods investigated in this work; specifically, we measure the overlap between the suggested MeSH terms and the MeSH terms included in the original query. We note that a MeSH term that is not in the original query may not necessarily be less effective of a search term than one included in the original query.

To evaluate the effectiveness of the suggested MeSH terms for the task of systematic review literature search, once query fragments are defragmented, the retrieval effectiveness is evaluated using typical systematic review literature search evaluation measures: precision, recall, and $F\beta$, with $\beta = \{1,3\}$. The PubMed Entrez API is used to directly issue defragmented Boolean queries to obtain retrieval results. As PubMed is constantly updated with new studies, we apply a date restriction to all queries for reproducibility purposes. We use the Jaccard index measure to evaluate the overlap of MeSH terms between those suggested by the investigated methods and those included in the original query.

For both evaluation settings (i.e., Boolean query retrieval and evaluation against original MeSH terms), we evaluate the lexical suggestion method in two settings: (i) **all**, where all retrieved MeSH terms are considered; and (i) **cut**, where the score-based cut-off is used. We also evaluate all BERTsuggestion methods and compare their effectiveness with that of the original query and the lexical methods.

2.5. Experimental setup

For our experiments, we use topics from the CLEF TAR task from 2017, 2018, and 2019 (Kanoulas et al., 2017, 2019, 2018). 15 topics are discarded due to lack of MeSH terms[5]. An additional topic is discarded because of retrieval issues[6], likely resulting from the fact that we translate queries automatically from one format (Ovid Medline) into another format (PubMed) (Scells et al., 2018). In total, we used 116 unique topics, as each year has partial overlap. For each topic, we automatically divide the Boolean query for that topic into query fragments (Scells et al., 2018). Each fragment contains at least one MeSH term. This results in a total of 311 unique query fragments (2.68 fragments per query on average). For each query fragment, we corrected any errors (e.g., spelling mistakes, syntactic errors), extracted MeSH terms, keywords, query fragments with MeSH terms, and query fragments without MeSH terms. For training the LTR model for each lexical

[5] Discarded topics are: **2017**: CD007427, CD010771, CD010772, CD010775, CD010783, CD010860, CD011145; **2018**: CD007427, CD009263, CD009694; **2019**: CD006715, CD007427, CD009263, CD009694, CD011768.

[6] The additional discard topic is **2017**: CD010276.

Table 2
Jaccard index(Jaccard) values quantifying the overlap between the MeSH terms suggsted by the investigated methods and those in the original query, along with the average number (Num) of MeSH term suggested by each method. In the original queries, there were on average 4.1343 MeSH terms for 2017, 4.8333 for 2018, 4.4000 for 2019-dta, and 2.7547 for 2019-intervention. Lexical methods: *CUT* indicates cut-off ranks. BERT methods: *FO, SA, SO, LN* indicate different cut-off strategies. Two-tailed statistical significance (*t*-test, $p < .05$) with Bonferroni correction between ATM and the other methods is indicated by $*$.

Dataset		2017		2018		2019-dta		2019-intervention	
Method		Jaccard	Num	Jaccard	Num	Jaccard	Num	Jaccard	Num
Lexical Method	ATM	0.0999	5.5373	0.2368	6.0139	0.2117	5.1500	0.2356	4.8868
	ATM-CUT	0.1995*	2.4179	0.1938	2.3056	0.2004	2.0500	0.2109	1.3019
	MetaMap	0.2654*	4.6866	0.2218	4.0417	0.2163	4.8000	0.2069	4.5094
	MetaMap-CUT	0.2374*	2.3134	0.1964	1.9028	0.2241	2.3500	0.1981	1.7736
	UMLS	0.2243*	8.9254	0.2235	7.9722	0.1905	7.7000	0.2405	7.5660
	UMLS-CUT	0.2751*	1.8955	0.2424	1.8611	0.1986	2.2000	0.2050	1.7547
	Fusion	0.2165*	11.4776	0.2160	10.9444	0.1735	10.5000	0.2212	9.7358
	Fusion-CUT	0.2761*	2.7761	0.2742	3.3194	0.2508	3.1000	0.2909	2.4340
BERT Method	Atomic-BERT-FO	0.2532*	12.7313	0.3105	12.2639	0.1573	11.8500	0.2252	13.6226
	Semantic-BERT-FO	0.2370*	11.0746	0.2963	10.6944	0.1654	10.7500	0.2219	11.5283
	Fragment-BERT-FO	0.3455*	1.0000	0.3812*	1.0000	0.1681	1.0000	0.2235	1.0000
	Fragment-BERT-SA	0.2233*	**16.6269**	0.2639	**16.4861**	0.1790	**15.5000**	0.2531	**17.2264**
	Fragment-BERT-SO	0.3921*	4.1343	0.4634*	4.8333	0.2574	4.4000	**0.3301**	2.7547
	Fragment-BERT-LN	0.2780*	5.2687	0.2689	3.7778	**0.2667**	3.8500	0.2415	3.8491

method, we use the pre-split training and test portions from the CLEF datasets. The 2019 topics are also split on systematic review type (intervention and diagnostic test accuracy — indicated as 'intervention' and 'dta' respectively in the results), while those for 2017 and 2018 are all diagnostic test accuracy. We use the 'quickrank' library (Capannini et al., 2016) for LTR, instantiated with LambdaMART trained to maximise nDCG. We leave other settings as per default.

For learning the linear function of the BERT suggestion method to decide the cut-off value of the MeSH term ranking list, as described in Section 2.3, we use the training portions of the CLEF TAR datasets. First, we obtain all the fragmentsfrom the CLEF TAR training splits. We count the number of free text atomic clausesand MeSH terms in each fragment. We then perform linear regression on these numbers to determine a function for each CLEF TAR dataset. We show the linear regression in Fig. 5.

3. Results

Results in this section are presented on the test splits of the CLEF TAR datasets (i.e., 2017, 2018, 2019-dta, 2019-intervention). We first analyse the search effectiveness of lexical methods versus our new BERT methods and then analyse the MeSH suggestion effectiveness compared to the MeSH terms originally used.

3.1. Retrieval effectiveness

Lexical Methods The results of the lexical methods presented in Table 4 are the same reported in our previous work Wang et al. (2021a). We discuss them briefly here for completeness. Unrefined methods generally contain higher recall than corresponding refined methods, with lower precision. This finding indicates that adding more MeSH terms in the query fragments can cause both more relevant and irrelevant studies to be retrieved. When compared using F1 and F3, the refined methods consistently outperform unrefined methods for each dataset. The U-CUT method achieves the highest effectiveness on CLEF 2017 and 2018, while A-CUT can achieve the highest effectiveness on CLEF 2019 dta; M-CUT for CLEF TAR 2019 intervention dataset in terms of F1 and F3. In terms of recall, the unrefined fusion method achieves the highest recall among all lexical suggestion methods. This gain in the recall is likely because unrefined fusion combines all of the MeSH terms suggested by the other three methods (ATM, MetaMap, and UMLS) using 'OR'. This suggests that the unrefined fusion method is not beneficial for improving the precision of a Boolean query. However, suppose semi-automatic MeSH term suggestion can be used. Information specialists may be able to use the suggestion and apply their expertise to decide which MeSH terms may be included to achieve higher performance.

BERT Methods

We first compare the effectiveness of the BERT methods with the original Boolean query. The result shows that under all evaluation measures (Precision, F1, F3 and Recall), BERT methods can outperform the original query on CLEF TAR 2017, 2018 and 2019-intervention, while effectiveness is generally worse for CLEF TAR 2019-dta. Note that CLEF TAR 2019-dta only contains eight unique topics; the lower effectiveness is likely due to a handful of topics.

Next, we compare the effectiveness of BERT suggestions against lexical suggestions. When comparing with un-refined lexical methods, the effectiveness of BERT suggestions is comparable in terms of F1 and F3, showing substantial gains across all datasets. However, compared with refined lexical methods, BERT suggestions generally obtain comparable results to refined lexical suggestions, except in CLEF TAR 2019-dta, in which refined lexical suggestion methods achieve higher effectiveness. In terms of recall, BERT suggestions obtain slightly higher recall to un-refined lexical suggestions, but substantially higher recall than refined lexical suggestions.

As mentioned in Section 3.1.0.1, unrefined lexical methods are effective to achieve higher recall while refined lexical methods are effective to achieve higher Precision, F1 and F3. We find that the MeSH terms suggested by BERT can obtain similar recall effectiveness to unrefined lexical methods while F1 and F3 can be comparable to refined lexical methods. Therefore, compare with lexical methods, BERT methods may be preferred to suggest more effective MeSH Terms.

3.2. Impact of BERT ranking representations

We compare different ranking representations of BERT, including Atomic BERT, Semantic BERTand Fragment BERT. We use the same cut-off strategy to compare these three representations fairly. We find that the precision, F1 and F3 values of Fragment BERTare the highest among the three methods, while recall of Fragment BERTis the lowest. However, only one MeSH term is suggested for each fragmentwhen Fragment BERTis used. This trade-off of precision and recall also suggests the same finding we described for lexical methods, where adding more MeSH terms can cause more studies to be retrieved. Between Semantic BERTand Atomic BERT, Semantic BERTis able to obtain higher precision while recall is lower than Atomic BERT. When comparing using F1 or F3, Semantic BERTalways achieves higher effectiveness. Therefore, the use of Semantic BERTis preferred over Atomic BERT.

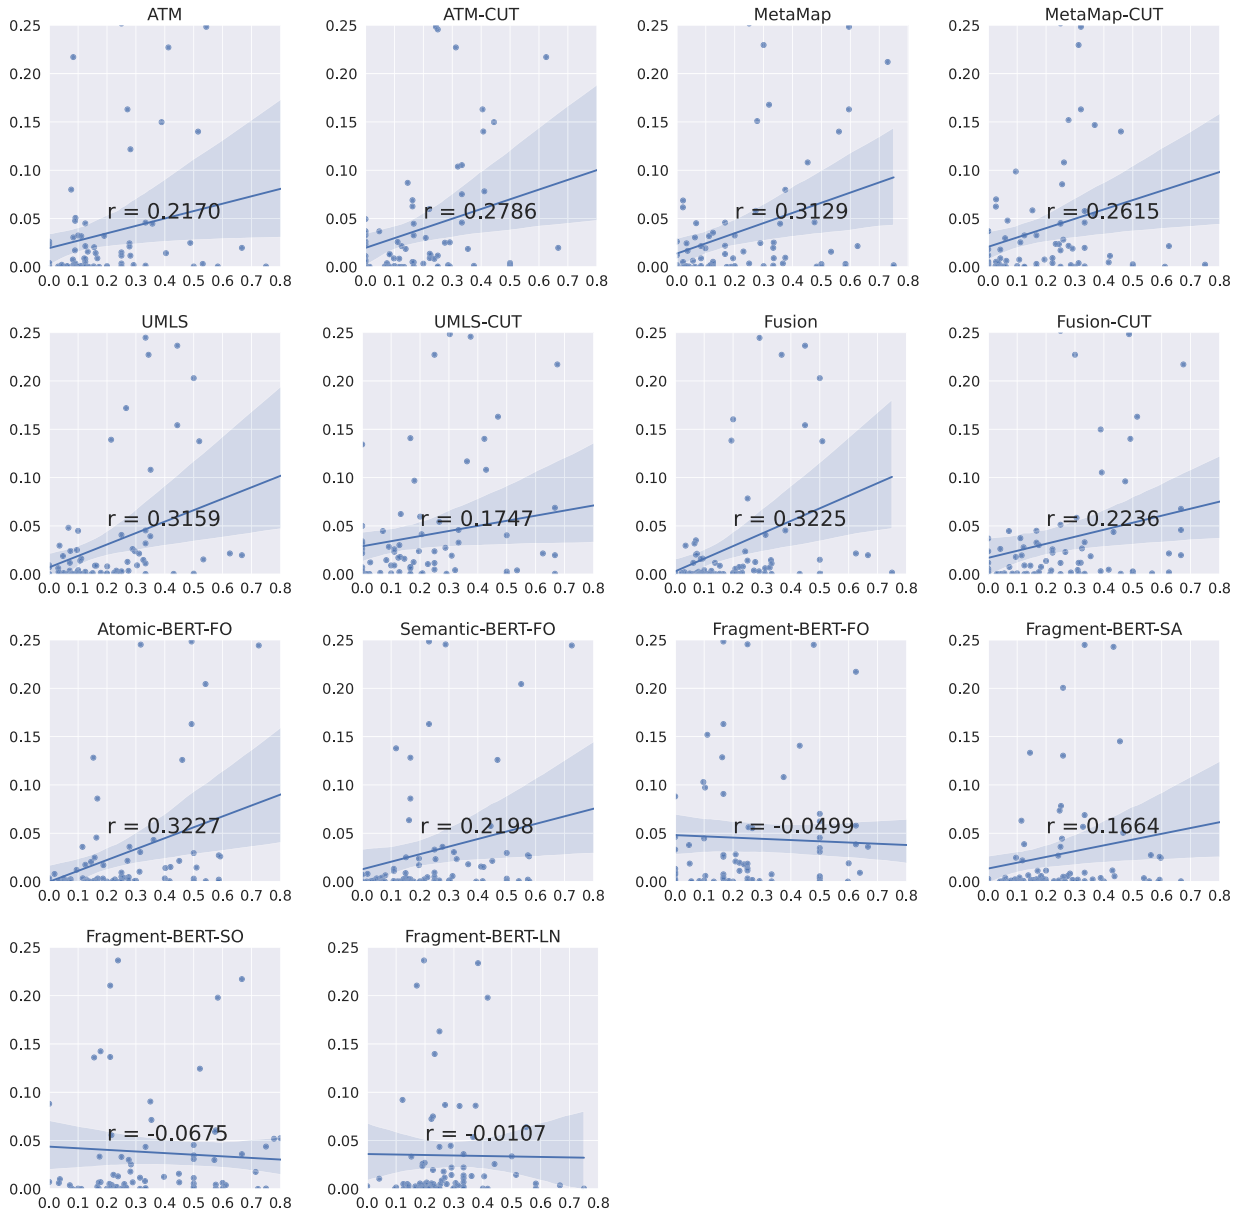

Fig. 6. Correlation graph of search effectiveness versus the overlap of MeSH terms. The x-axis reports the Jaccard index for overlap between suggested MeSH terms and MeSH terms included in the original query, and the y-axis reports the F1 values for search effectiveness for each topic.

3.3. Impact of cut-off strategy

When comparing different cut-off strategies for BERT suggestions[7], we find that FO can consistently achieve the highest Precision, F1 and F3 compared to other cut-off methods. On the other hand, the recall value of using FO is the lowest among all other methods, indicating that the trade-off of precision and recall is again caused by the number of MeSH terms added to the query. For the other three cut-off strategies, including SA, SO, and LN, we find that SO and LN consistently outperform SA, suggesting that information specialists have an intuition for how many MeSH terms to add to a query.

[7] Only Fragment BERTis considered as SA, SO, LN are only applicable to Fragment BERT.

3.4. Are suggested MeSH terms the same as those in the original queries?

Next, we study the overlap between the MeSH term suggested by the considered methods and those included in the original query; this is reported in Table 2 and is measured with the Jaccard index. One immediate observation is that the overlap of all based methods is considerably higher than that of lexical methods. This observation is based on that the highest value of the Jaccard index in each dataset always appear in the BERT suggestion method. Moreover, when applying the SO cut-off strategy to Fragment BERT, the highest overlap is always obtained, which indicates that BERT suggestion methods also agree on the Terms chosen by systematic reviewers..

The previous results reported in Table 4 highlighted that, in general, BERT methods were better than lexical methods in suggesting effective search terms; and these were more effective than those in the original queries, although differences were not statistically significant. These results, in conjunction with the findings in Table 2, indicate that the BERT methods identify very similar MeSH terms that present in the

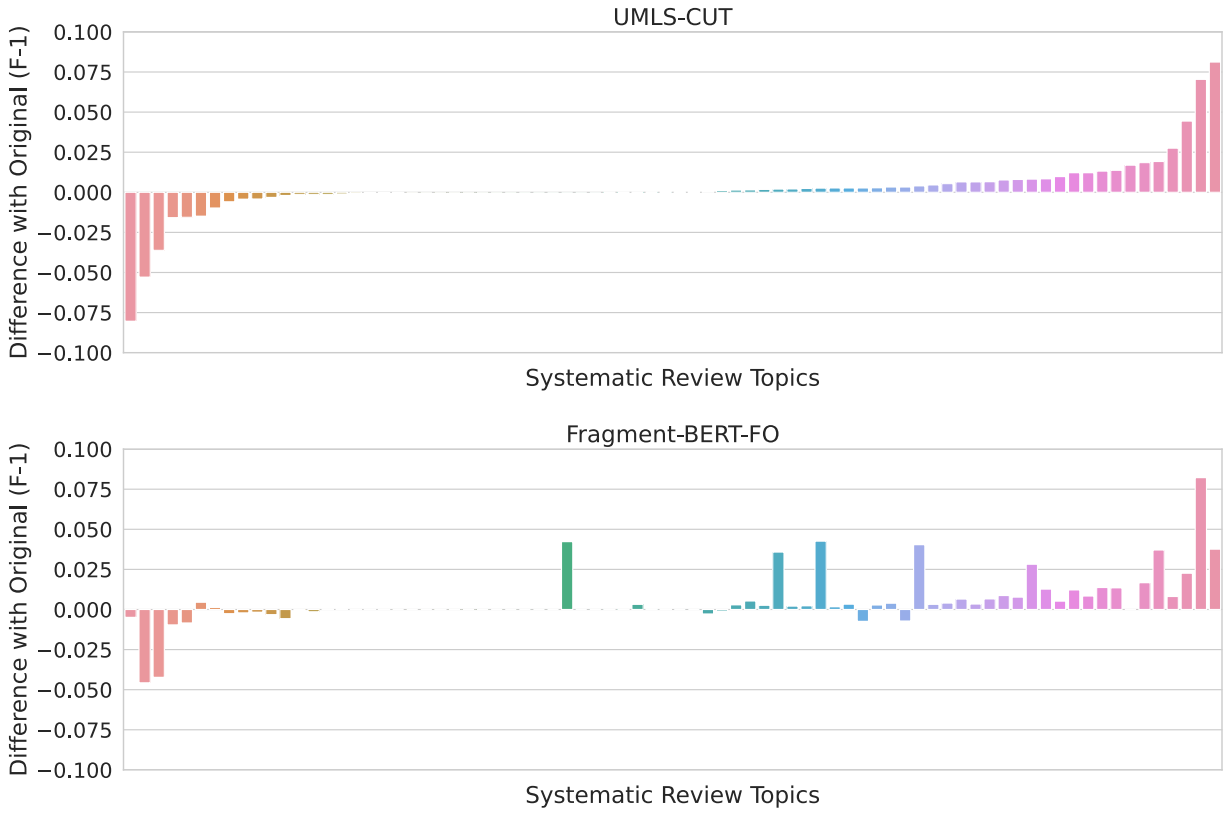

Fig. 7. Plot showing systematic review topics versus original query effectiveness; each bar represents a topic. The y-axis represents the effectiveness difference between the query with the suggested MeSH terms and the original query. Effectiveness is measured using F1.

original queries – and the MeSH terms identified by BERT methods are more effective than those provided by other methods.

Intuitively, we further analyse whether search effectiveness and the suggestion of MeSH terms that are included in the original query correlate, meaning that Mesh Terms used in the original Boolean query may be of very high quality and should be used as gold standard. The Jaccard index measure is used once more to represent the similarity between suggested MeSH terms and those present in the original query, and F1 is used to represent the search effectiveness of the associated query (with suggested MeSH terms included). The results of this correlation analysis are reported in Fig. 6. We find that, while for all lexical methods search effectiveness is weakly correlated with the overlap of MeSH terms, this is not the case for BERT methods. This indicates that MeSH term suggestionfrom the original query may not be the best MeSH term suggestionto suggest. In fact, it often is that MeSH terms that are suggested but not included in the original query provide higher search effectiveness than the original MeSH terms themselves.

3.5. Search stability

We next analyse the search effectiveness stability of different MeSH term suggestion methods on a topic-by-topic basis. With search effectiveness stability we refer to the amount of variance across topics of the measured search effectiveness obtained when using queries with MeSH

Table 3
Search effectiveness of the Boolean queries with the suggested MeSH terms evaluated by precision (P), F1, F3 and recall (R). For Lexical methods: *CUT* indicates cut-off ranks. BERT methods: *FO, SA, SO, LN* indicate different cut-off strategies.

Topic ID	CD009642				CD004414			
Method	P	F1	F3	R	P	F1	F3	R
ORIGINAL	0.0088	0.0175	0.0344	1.0000	0.0013	0.0026	0.0052	0.6875
ATM	0.0109	0.0215	0.0421	0.9194	0.0018	0.0035	0.0070	0.3125
ATM-CUT	0.0109	0.0215	0.0421	0.9194	0.0020	0.0040	0.0078	0.3125
MetaMap	0.0109	0.0215	0.0421	0.9194	0.0018	0.0035	0.0070	0.3125
MetaMap-CUT	0.0109	0.0215	0.0421	0.9194	0.0014	0.0027	0.0054	0.3125
UMLS	0.0109	0.0215	0.0421	0.9194	0.0013	0.0025	0.0050	0.3125
UMLS-CUT	0.0109	0.0215	0.0421	0.9194	0.0020	0.0040	0.0078	0.3125
Fusion	0.0109	0.0215	0.0421	0.9194	0.0018	0.0035	0.0069	0.3125
Fusion-CUT	0.0109	0.0215	0.0421	0.9194	0.0014	0.0027	0.0054	0.3125
Atomic-BERT-FO	0.0108	0.0214	0.0418	0.9194	0.0012	0.0024	0.0048	0.3125
Semantic-BERT-FO	0.0108	0.0214	0.0418	0.9194	0.0012	0.0024	0.0048	0.3125
Fragment-BERT-SA	0.0259	0.0504	0.0955	0.9194	0.0012	0.0024	0.0048	0.3125
Fragment-BERT-SO	0.0270	0.0525	0.0993	0.9194	0.0028	0.0055	0.0109	0.3125
Fragment-BERT-LN	0.0276	0.0536	0.1013	0.9194	0.0013	0.0026	0.0052	0.3125

terms suggested by a specific MeSH term suggestion method. The larger the effectiveness, the lower the stability. We only analyse the best-performing lexical (U-CUT) and BERT (F-B-FO) methods.

Fig. 7, which combines the topics of all of the CLEF TAR datasets test splits, shows that for most of the topics, both kinds of MeSH suggestion methods can outperform or match the effectiveness of the original queries. We also find that our MeSH term suggestion methods sometimes obtain lower effectiveness. It is unclear if these are difficult topics to suggest MeSH terms for, or if there are mistakes in these queries that cause the poor effectiveness (e.g., spelling mistakes in the free text atomic clausesthat were not detected at the time of data cleaning).

3.6. Case study

Given the findings above, we next seek to investigate the reasons for highly effective or ineffective results. We choose topic CD009642 and topic CD004414 from the CLEF TAR 2019-intervention dataset as they are representative topics where suggestion methods outperform the effectiveness (CD009642) or struggle to match the effectiveness (CD004414) of the original query. Query fragments corresponding to these topics for all the suggestion methods are shown in Tables 6 and 7. We also show their search effectiveness in Table 3.

Firstly, we find that both the suggested MeSH terms and the search effectiveness are similar for all lexical methods. One exception is UMLS in topic CD004414, which suggests more MeSH terms than the other two methods, causing a drop in effectiveness.

On the other hand, MeSH terms suggested by BERT methods appear to differ greatly from lexical methods. BERT methods have captured both lexically similar MeSH terms and terms semantically related to the input free text atomic clauses. One example is shown in the first fragment of topic CD009642. While all lexical methods suggest *Lidocaine* which is lexically equal to *lidocain*, BERT methods suggest similar drugs such as *Procaine*. Another example in topic CD004414 shows that BERT methods can use this semantic matching ability to suggest MeSH terms indicating the method of intervention, shown in suggesting *Patch Tests*. Therefore, BERT methods suggest MeSH terms that are not bound to the lexical semantics of a free text atomic clause. Another advantage of BERT methods is that they guarantee that at least one MeSH term will be suggested. For lexical methods, suggestions are based on pre-existing rule-based knowledge; thus, when free text atomic clausescan not be matched, no MeSH terms can be suggested (e.g., ATM does not suggest any MeSH term in Fragment 1 of CD009642). Overall, we believe that the semantic matching of BERT may sometimes be detrimental to MeSH suggestion. We leave the investigation into how to prevent BERT from suggesting MeSH terms that are not relevant to the information need of a query fragment (e.g., suggesting a MeSH term that is the intervention of an outcome for a query fragment) for future work.

4. Conclusion

In this article, we extend our previous line of research on suggesting MeSH terms for Boolean queries for systematic review literature search. This task adds to a recent stream of research that has focused on computational methods for the assisted creation (Agosti et al., 2019; Scells et al., 2020a, 2021, 2020b) or refinement (Agosti et al., 2020; Alharbi & Stevenson, 2020; Harrisen & Guido, 2018; Scells et al., 2019; Wang et al., 2021a) of Boolean queries for systematic review creation. In addition to the lexical methods we proposed previously, in this new line of work, we introduced a new set of BERT based MeSH suggestion methods. We undertook a comprehensive evaluation and analysis of our new MeSH suggestion methods. We compared the effectiveness of the suggested MeSH terms from our new methods to both our existing lexical methods and the original queries formulated by information specialists. We found that the MeSH terms originally chosen by information specialists were often not the most effective choice and that more effective MeSH terms can be suggested automatically by our new methods. We also found that using BERT methods can generally achieve higher effectiveness than the Lexical method in MeSH Term suggestion: this may be due to the fact that BERT methods were often able to capture deeper semantic relationships. This finding motivates future work to combine lexical and BERT methods in order to reap the benefits of both approaches. Combining such sparse and dense approaches has seen much success in related areas of research, such as ad-hoc search (Karpukhin et al., 2020; Li et al., 2022; Ma et al., 2021; Wang et al., 2021b).

In addition, we believe that the full potential of using MeSH entities in our suggestion method is unexplored. In our future work, we project three research directions using more information from MeSH entities to achieve more effective MeSH term suggestions, including (1) Use of MeSH tree hierarchy: MeSH entities are organised in a tree hierarchy. The parent-child relationship of entities may further restrict the number of MeSH terms suggested by MeSH Term suggestion methods (Exp: Use parent MeSH entity to restrict which child entities can appear in the suggestion list). (2) Use of MeSH categories: MeSH entities are categorised according to their natures (term, concept, descriptor and category). The nature of MeSH entities may be used in the fine-tuning process of the MeSH term suggestion methods to represent the MeSH entities. (3) Use of external MeSH definition: Each MeSH entity has a corresponding Wikipedia page to explain its content and uses. These comprehensive pages may be used to further fine-tune our MeSH term suggestion model to achieve effective MeSH term suggestions.

Identifying MeSH terms to add to a Boolean query for a systematic review literature search is a difficult task for information specialists. The findings of this article have implications for both the Information Retrieval and Systematic Review communities. Firstly, our methods can be used in automatic query formulation situations (see, e.g., work by Scells et al. (2021)). Secondly, they can be integrated into existing tools to assist information specialists in formulating more effective queries (Li et al., 2020; Scells & Zuccon, 2018).

CRediT authorship contribution statement

Shuai Wang: Conceptualization, Methodology, Software, Investigation, Formal analysis, Writing – original draft. **Harrisen Scells:** Writing – review & editing, Formal analysis, Validation, Methodology. **Bevan Koopman:** Supervision, Conceptualization. **Guido Zuccon:** Supervision, Conceptualization, Methodology, Funding acquisition, Project administration, Writing – review & editing.

Declaration of Competing Interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Acknowledgments

Shuai Wang is supported by a UQ Earmarked PhD Scholarship and this research is funded by the Australian Research Council Discovery Projects program ARC Discovery Project DP210104043.

Appendices

Table 4
Search effectiveness of Boolean query using suggested MeSH terms evaluated by precision (P), F1, F3 and recall (R). Lexical methods: For each method, *CUT* indicates cut-off ranks. BERT methods: *FO, SA, SO, LN* indicate different cut-off strategies. No statistical significant differences are detected between the ORIGINAL query and those obtained by the other methods (two-tailed *t*-test with Bonferroni correction, $p < .05$).

Dataset		2017				2018				2019-dta				2019-intervention			
Method		P	F1	F3	R	P	F1	F3	R	P	F1	F3	R	P	F1	F3	R
	ORIGINAL	0.0288	0.0311	0.0440	0.7745	0.0323	0.0576	0.0965	0.8629	0.0227	**0.0421**	**0.0738**	0.8966	0.0165	0.0212	0.0309	0.7471
Lexical Method	ATM	0.0265	0.0262	0.0353	0.7549	0.0317	0.0552	0.0898	0.8190	0.0113	0.0211	0.0373	0.8916	0.0156	0.0183	0.0269	0.7073
	ATM-CUT	0.0316	0.0299	0.0404	0.7269	0.0354	0.0624	0.1033	0.7998	**0.0243**	0.0398	0.0637	0.8375	0.0173	0.0191	0.0288	0.6938
	MetaMap	0.0304	0.0287	0.0381	0.7519	0.0342	0.0599	0.0980	0.8150	0.0131	0.0245	0.0433	0.8791	0.0135	0.0218	0.0339	0.6974
	MetaMap-CUT	0.0337	0.0312	0.0423	0.7191	0.0360	0.0633	0.1043	0.8071	0.0193	0.0358	0.0625	0.8393	0.0159	0.0251	0.0382	0.6831
	UMLS	0.0275	0.0269	0.0355	0.7458	0.0297	0.0519	0.0847	0.8200	0.0114	0.0214	0.0384	0.8616	0.0118	0.0183	0.0275	0.6998
	UMLS-CUT	0.0335	0.0315	0.0430	0.7225	0.0384	**0.0681**	**0.1133**	0.7963	0.0174	0.0305	0.0508	0.8381	0.0173	0.0191	0.0295	0.6638
	Fusion	0.0218	0.0227	0.0300	0.7712	0.0284	0.0495	0.0800	0.8455	0.0103	0.0192	0.0342	0.9075	0.0109	0.0173	0.0263	0.7212
	Fusion-CUT	0.0323	0.0303	0.0409	0.7282	0.0333	0.0582	0.0951	0.8120	0.0147	0.0269	0.0465	0.8394	0.0161	0.0173	0.0262	0.6797
BERT Method	Atomic-BERT-FO	0.0257	0.0249	0.0330	**0.7830**	0.0289	0.0488	0.0795	0.8523	0.0092	0.0173	0.0310	0.8870	0.0070	0.0126	0.0219	0.7587
	Semantic-BERT-FO	0.0273	0.0272	0.0363	0.7633	0.0284	0.0501	0.0820	0.8502	0.0096	0.0181	0.0324	0.8870	0.0110	0.0183	0.0288	0.7483
	Fragment-BERT-FO	**0.0342**	**0.0324**	**0.0446**	0.7415	**0.0382**	0.0678	0.1132	0.8041	0.0169	0.0314	0.0548	0.8924	**0.0212**	**0.0276**	**0.0422**	0.7106
	Fragment-BERT-SA	0.0212	0.0216	0.0284	0.7699	0.0268	0.0471	0.0772	**0.8652**	0.0097	0.0181	0.0323	**0.9357**	0.0076	0.0137	0.0235	**0.7806**
	Fragment-BERT-SO	0.0265	0.0250	0.0335	0.7593	0.0328	0.0588	0.0991	0.8258	0.0129	0.0243	0.0433	0.8987	0.0176	0.0238	0.0358	0.7431
	Fragment-BERT-LN	0.0265	0.0274	0.0373	0.7615	0.0318	0.0561	0.0925	0.8355	0.0112	0.0211	0.0378	0.8969	0.0105	0.0167	0.0265	0.7428

Table 5
Two-tailed *t*-test results of Boolean query search effectiveness between the ORIGINAL query and those obtained by the other methods by precision (P), F1, F3 and recall (R). Lexical methods: *CUT* indicates cut-off ranks. BERT methods: *FO, SA, SO, LN* indicate different cut-off strategies.

Dataset		2017				2018				2019-dta				2019-intervention			
Method		P	F1	F3	R	P	F1	F3	R	P	F1	F3	R	P	F1	F3	R
Lexical Method	ATM	0.9148	0.7515	0.6465	0.8570	0.9670	0.9216	0.8578	0.5820	0.4586	0.4573	0.4544	0.9674	0.9400	0.7730	0.7292	0.7467
	ATM-CUT	0.8992	0.9396	0.8480	0.6698	0.8426	0.8532	0.4265	0.9196	0.9378	0.8354	0.6629	0.9514	0.8275	0.8556	0.6688	0.8306
	MetaMap	0.9411	0.8784	0.7572	0.8366	0.9263	0.9668	0.5474	0.5297	0.5302	0.5294	0.8930	0.7702	0.9612	0.8361	0.6867	0.8956
	MetaMap-CUT	0.8279	0.9954	0.9320	0.6203	0.8159	0.8330	0.4813	0.8305	0.8271	0.8205	0.6691	0.9502	0.7414	0.6170	0.6078	0.7990
	UMLS	0.9468	0.7873	0.6529	0.7933	0.8132	0.7511	0.5876	0.4432	0.4475	0.4535	0.7712	0.6556	0.8076	0.8162	0.7007	0.8556
	UMLS-CUT	0.8323	0.9785	0.9583	0.6363	0.6625	0.6500	0.4027	0.7251	0.6698	0.6245	0.6667	0.9508	0.8238	0.8980	0.5085	0.6661
	Fusion	0.7208	0.5824	0.4505	0.9755	0.7373	0.6598	0.8258	0.4014	0.4031	0.4049	0.9199	0.5801	0.7308	0.7487	0.8293	0.7838
	Fusion-CUT	0.8768	0.9598	0.8729	0.6786	0.9817	0.9714	0.5223	0.5972	0.5856	0.5704	0.6722	0.9726	0.6744	0.6682	0.5919	0.9463
BERT Method	Atomic-BERT-FO	0.8905	0.7021	0.5662	0.9363	0.8109	0.7149	0.6493	0.8917	0.3493	0.3526	0.3566	0.9410	0.2777	0.3441	0.4487	0.9215
	Semantic-BERT-FO	0.9465	0.8073	0.6916	0.9164	0.7530	0.6962	0.8704	0.3615	0.3657	0.3711	0.9410	0.5684	0.7928	0.8853	0.9918	0.7841
	Fragment-BERT-FO	0.8075	0.9371	0.9712	0.7663	0.6737	0.6540	0.4589	0.7074	0.7015	0.6948	0.9715	0.7003	0.5612	0.4209	0.7636	0.6814
	Fragment-BERT-SA	0.7171	0.5460	0.3968	0.9658	0.6480	0.5888	0.9754	0.3723	0.3748	0.3777	0.6593	0.3085	0.4157	0.5481	0.7687	0.6837
	Fragment-BERT-SO	0.9135	0.6889	0.5615	0.8908	0.9592	0.9406	0.6374	0.5109	0.5154	0.5208	0.9856	0.9295	0.8150	0.7191	0.9736	0.9739
	Fragment-BERT-LN	0.9119	0.8152	0.7293	0.9053	0.9452	0.9110	0.7120	0.4368	0.4416	0.4474	0.9978	0.5192	0.6417	0.7200	0.9715	0.9681

Table 6
Query fragments in different methods, For Lexical methods: *CUT* indicates cut-off ranks. For BERT suggestion method, *A-B*indicates Atomic BERT, *S-B* indicates Semantic BERT, *F-B* indicates Fragment BERT. In each BERT method, *FO, SA, SO, LN* indicates cut-off strategy used. For each fragment, bold text means MeSH term.

Topic	CD009642		
Fragments	Fragment 1		Fragment 2
ORIGINAL	**Lidocaine** OR lidocain* OR Lignocain* OR Xylocain*		**Pain** OR **Pain, Postoperative** OR **Postoperative Care** OR **Postoperative Complications** OR (post operative OR postoperative) AND (pain* OR recovery)
ATM	lidocain* OR Lignocain* OR Xylocain*		**Pain** OR (post operative OR postoperative) AND (pain* OR recovery)
ATM-CUT	lidocain* OR Lignocain* OR Xylocain*		**Pain** OR (post operative OR postoperative) AND (pain* OR recovery)
MetaMap	**Lidocaine** OR lidocain* OR Lignocain* OR Xylocain*		**Pain** OR (post operative OR postoperative) AND (pain* OR recovery)
MetaMap-CUT	**Lidocaine** OR lidocain* OR Lignocain* OR Xylocain*		**Pain** OR (post operative OR postoperative) AND (pain* OR recovery)
UMLS	**Lidocaine** OR lidocain* OR Lignocain* OR Xylocain*		**Pain** OR (post operative OR postoperative) AND (pain* OR recovery)
UMLS-CUT	**Lidocaine** OR lidocain* OR Lignocain* OR Xylocain*		**Pain** OR (post operative OR postoperative) AND (pain* OR recovery)
Fusion	**Lidocaine** OR lidocain* OR Lignocain* OR Xylocain*	AND	**Pain** OR (post operative OR postoperative) AND (pain* OR recovery)
Fusion-CUT	**Lidocaine** OR lidocain* OR Lignocain* OR Xylocain*		**Pain** OR (post operative OR postoperative) AND (pain* OR recovery)
Atomic-BERT-FO	**Lidocaine** OR **Xylans** OR lidocain* OR Lignocain* OR Xylocain*		**Postoperative Care** OR **Pain** OR **Recovery of Function** OR (post operative OR postoperative) AND (pain* OR recovery)
Semantic-BERT-FO	**Lidocaine** OR **Xylans** OR lidocain* OR Lignocain* OR Xylocain*		**Postoperative Care** OR **Pain** OR **Recovery of Function** OR (post operative OR postoperative) AND (pain* OR recovery)
Fragment-BERT-FO	**Lidocaine** OR lidocain* OR Lignocain* OR Xylocain*		**Pain, Postoperative** OR (post operative OR postoperative) AND (pain* OR recovery)
Fragment-BERT-SA	**Lidocaine** OR **Procaine** OR **Xylans** OR lidocain* OR Lignocain* OR Xylocain*		**Pain, Postoperative** OR **Postoperative Care** OR **Postoperative Period** OR **Postoperative Complications** OR (post operative OR postoperative) AND (pain* OR recovery)
Fragment-BERT-SO	**Lidocaine** OR lidocain* OR Lignocain* OR Xylocain*		**Pain, Postoperative** OR **Postoperative Care** OR **Postoperative Period** OR **Postoperative Complications** OR (post operative OR postoperative) AND (pain* OR recovery)
Fragment-BERT-LN	**Lidocaine** OR **Procaine** OR **Xylans** OR lidocain* OR Lignocain* OR Xylocain*		**Pain, Postoperative** OR **Postoperative Care** OR **Postoperative Period** OR (post operative OR postoperative) AND (pain* OR recovery)

Table 7
Query fragments in different methods, For Lexical methods: *CUT* indicates cut-off ranks. For BERT suggestion method, *A-B*indicates Atomic BERT, *S-B* indicates Semantic BERT, *F-B* indicates Fragment BERT. In each BERT method, *FO, SA, SO, LN* indicates cut-off strategy used. For each fragment, bold text means MeSH term.

Topic	CD004414		
Fragments	Fragment 1		Fragment 2
ORIGINAL	**Hand** OR hand* OR finger* OR palm*		**Hand Dermatoses** OR (dermat* OR eczema) AND (occupation* OR irritant* OR contact) AND (hand* OR finger* OR palm*)
ATM	**Hand** OR **Fingers** OR hand* OR finger* OR palm*		**Fingers** OR **Eczema** OR **Hand** OR **Irritants** OR **Occupations** OR (dermat* OR eczema) AND (occupation* OR irritant* OR contact) AND (hand* OR finger* OR palm*)
ATM-CUT	**Hand** OR hand* OR finger* OR palm*		**Fingers** OR **Eczema** OR (dermat* OR eczema) AND (occupation* OR irritant* OR contact) AND (hand* OR finger* OR palm*)
MetaMap	**Hand** OR **Fingers** OR hand* OR finger* OR palm*		**Hand** OR **Fingers** OR **Eczema** OR **Occupations** OR **Irritants** OR (dermat* OR eczema) AND (occupation* OR irritant* OR contact) AND (hand* OR finger* OR palm*)
MetaMap-CUT	**Hand** OR hand* OR finger* OR palm*		**Hand** OR **Fingers** OR **Eczema** OR (dermat* OR eczema) AND (occupation* OR irritant* OR contact) AND (hand* OR finger* OR palm*)
UMLS	**Fingers** OR **Hand** OR **Palm Oil** OR **Computers, Handheld** OR hand* OR finger* OR palm*		**Eczema** OR **Fingers** OR **Hand** OR **Dermatitis, Atopic** OR **Kaposi Varicelliform Eruption** OR **Retirement** OR **Computers, Handheld** OR **Occupations** OR **Palm Oil** OR **Irritants** OR (dermat* OR eczema) AND (occupation* OR irritant* OR contact) AND (hand* OR finger* OR palm*)
UMLS-CUT	**Fingers** OR hand* OR finger* OR palm*		**Eczema** OR **Fingers** OR (dermat* OR eczema) AND (occupation* OR irritant* OR contact) AND (hand* OR finger* OR palm*)
Fusion	**Hand** OR **Fingers** OR **Palm Oil** OR **Computers, Handheld** OR hand* OR finger* OR palm*	AND	**Eczema** OR **Fingers** OR **Hand** OR **Dermatitis, Atopic** OR **Occupations** OR **Kaposi Varicelliform Eruption** OR **Retirement** OR **Irritants** OR **Computers, Handheld** OR **Palm Oil** OR (dermat* OR eczema) AND (occupation* OR irritant* OR contact) AND (hand* OR finger* OR palm*)
Fusion-CUT	**Hand** OR hand* OR finger* OR palm*		**Eczema** OR **Fingers** OR **Hand** OR (dermat* OR eczema) AND (occupation* OR irritant* OR contact) AND (hand* OR finger* OR palm*)
Atomic-BERT-FO	**Hand** OR **Fingers** OR **Palm Oil** OR hand* OR finger* OR palm*		**Hand** OR **Eczema** OR **Occupations** OR **Dermatology** OR **Irritants** OR **Dermatitis, Contact** OR

(*continued on next page*)

Table 7 (*continued*)

Topic	CD004414	
Fragments	Fragment 1	Fragment 2
		Fingers OR **Palm Oil** OR (dermat* OR eczema) AND (occupation* OR irritant* OR contact) AND (hand* OR finger* OR palm*)
Semantic-BERT-FO	**Hand** OR **Fingers** OR **Palm Oil** OR hand* OR finger* OR palm*	**Dermatology** OR **Eczema** OR **Occupations** OR **Irritants** OR **Dermatitis, Contact** OR **Hand** OR **Fingers** OR **Palm Oil** OR (dermat* OR eczema) AND (occupation* OR irritant* OR contact) AND (hand* OR finger* OR palm*)
Fragment-BERT-FO	**Hand** OR hand* OR finger* OR palm*	**Dermatitis, Contact** OR (dermat* OR eczema) AND (occupation* OR irritant* OR contact) AND (hand* OR finger* OR palm*)
Fragment-BERT-SA	**Hand** OR **Fingers** OR **Palm Oil** OR hand* OR finger* OR palm*	**Dermatitis, Contact** OR **Dermatitis, Allergic Contact** OR **Hand** OR **Fingers** OR **Eczema** OR **Dermatitis, Atopic** OR **Patch Tests** OR **Skin Diseases** OR (dermat* OR eczema) AND (occupation* OR irritant* OR contact) AND (hand* OR finger* OR palm*)
Fragment-BERT-SO	**Hand** OR hand* OR finger* OR palm*	**Dermatitis, Contact** OR (dermat* OR eczema) AND (occupation* OR irritant* OR contact) AND (hand* OR finger* OR palm*)
Fragment-BERT-LN	**Hand** OR **Fingers** OR **Palm Oil** OR hand* OR finger* OR palm*	**Dermatitis, Contact** OR **Dermatitis, Allergic Contact** OR **Hand** OR (dermat* OR eczema) AND (occupation* OR irritant* OR contact) AND (hand* OR finger* OR palm*)

References

Abdou, S., & Savoy, J. (2008). Searching in medline: Query expansion and manual indexing evaluation. *Information Processing & Management, 44*(2), 781–789.

Adlassnig, K., et al. (2009). Optimization of the pubmed automatic term mapping, *vol. 150. Medical informatics in a united and healthy europe: Proceedings of mie 2009, the xxii international congress of the european federation for medical informatics* (p. 238). IOS Press.

Agosti, M., Di Nunzio, G. M., & Marchesin, S. (2019). An analysis of query reformulation techniques for precision medicine. *Proceedings of the 42nd international acm sigir conference on research and development in information retrieval* (pp. 973–976).

Agosti, M., Di Nunzio, G. M., & Marchesin, S. (2020). *A post-analysis of query reformulation methods for clinical trials retrieval*.

Alharbi, A., & Stevenson, M. (2020). Refining boolean queries to identify relevant studies for systematic review updates. *Journal of the American Medical Informatics Association, 27*(11), 1658–1666.

Aronson, A. R. (2001). Effective mapping of biomedical text to the umls metathesaurus: the metamap program. *Proceedings of the amia symposium* (p. 17). American Medical Informatics Association.

Balog, K. (2018). *Entity-oriented search*. Springer.

Bird, S., & Loper, E. (2004). *Nltk: the natural language toolkit*. Association for Computational Linguistics.

Bodenreider, O. (2004). The unified medical language system (umls): integrating biomedical terminology. *Nucleic Acids Research, 32*(suppl_1), D267–D270.

Brown, T., Mann, B., Ryder, N., Subbiah, M., Kaplan, J. D., Dhariwal, P., Neelakantan, A., Shyam, P., Sastry, G., Askell, A., et al. (2020). Language models are few-shot learners. *Advances in Neural Information Processing Systems, 33*, 1877–1901.

Capannini, G., Lucchese, C., Nardini, F. M., Orlando, S., Perego, R., & Tonellotto, N. (2016). Quality versus efficiency in document scoring with learning-to-rank models. *Information Processing & Management, 52*(6), 1161–1177.

Chalkidis, I., Fergadiotis, M., Malakasiotis, P., Aletras, N., & Androutsopoulos, I. (2020). LEGAL-BERT: The muppets straight out of law school. *Findings of the association for computational linguistics: Emnlp 2020* (pp. 2898–2904). Online: Association for Computational Linguistics. https://doi.org/10.18653/v1/2020.findings-emnlp.261

Chang, A. A., Heskett, K. M., & Davidson, T. M. (2006). Searching the literature using medical subject headings versus text word with pubmed. *The Laryngoscope, 116*(2), 336–340.

Choe, S., Aum, S., & Kim, J. H. (2022). Short review on srbert: Automatic article classification model for systematic review using bert. *Asian Journal of Complementary and Alternative Medicine*, 16.

Clark, J. (2013). Systematic reviewing. In G. M. W. Suhail A. R. Doi (Ed.), *Methods of clinical epidemiology*. Springer.

Cormack, G. V., & Grossman, M. R. (2017). Technology-assisted review in empirical medicine: Waterloo participation in clef ehealth 2017. *Clef (working notes)*.

Devlin, J., Chang, M., Lee, K., & Toutanova, K. (2019). BERT: pre-training of deep bidirectional transformers for language understanding. In J. Burstein, C. Doran, & T. Solorio (Eds.), *Proceedings of the 2019 conference of the north american chapter of the association for computational linguistics: Human language technologies, NAACL-HLT 2019, minneapolis, mn, usa, june 2–7, 2019, volume 1 (long and short papers)* (pp. 4171–4186). Association for Computational Linguistics. https://doi.org/10.18653/v1/n19-1423.

Fox, E. A., & Shaw, J. A. (1994). Combination of multiple searches. *NIST special publication SP, 243*.

Gao, L., Ma, X., Lin, J. J., & Callan, J. (2022). Tevatron: An efficient and flexible toolkit for dense retrieval. *ArXiv, abs/2203.05765*.

Harrisen, S., & Guido, Z. (2018). Generating better queries for systematic reviews, . In *SIGIR '18The 41st international acm sigir conference on research & development in information retrieval* (pp. 475–484). New York, NY, USA: ACM.

Jimmy, Zuccon, G., Koopman, B., & Demartini, G. (2019). Health card retrieval for consumer health search: An empirical investigation of methods, . In *CIKM '19Proceedings of the 28th acm international conference on information and knowledge management* (pp. 2405–2408). New York, NY, USA: Association for Computing Machinery. https://doi.org/10.1145/3357384.3358128

Kanoulas, E., Li, D., Azzopardi, L., & Spijker, R. (2017). Clef 2017 technologically assisted reviews in empirical medicine overview, *vol. 1866. Ceur workshop proceedings* (pp. 1–29).

Kanoulas, E., Li, D., Azzopardi, L., & Spijker, R. (2019). Clef 2019 technology assisted reviews in empirical medicine overview, *vol. 2380. Ceur workshop proceedings*.

Kanoulas, E., Spijker, R., Li, D., & Azzopardi, L. (2018). Clef 2018 technology assisted reviews in empirical medicine overview. *Clef 2018 evaluation labs and workshop: Online working notes, ceur-ws*.

Karpukhin, V., Oğuz, B., Min, S., Lewis, P., Wu, L., Edunov, S., Chen, D., & Yih, W.-t. (2020). Dense passage retrieval for open-domain question answering. *arXiv preprint arXiv:2004.04906*.

Khosrovian, K., Pfahl, D., & Garousi, V. (2008). Gensim 2.0: A customizable process simulation model for software process evaluation. *International conference on software process* (pp. 294–306). Springer.

Lee, G. E., & Sun, A. (2018). Seed-driven document ranking for systematic reviews in evidence-based medicine, . In *SIGIR '18The 41st international acm sigir conference on research & development in information retrieval. (2018)* (pp. 455–464). New York, NY, USA: ACM. https://doi.org/10.1145/3209978.3209994. http://doi.acm.org/10.1145/3209978.3209994

Lee, G. E., & Sun, A. (2022). Towards reducing manual workload in technology-assisted reviews: Estimating ranking performance. *arXiv preprint arXiv:2201.05648*.

Lee, J., Yoon, W., Kim, S., Kim, D., Kim, S., So, C. H., & Kang, J. (2020). Biobert: A pre-trained biomedical language representation model for biomedical text mining. *Bioinformatics (Oxford, England), 36*(4), 1234–1240.

Li, D., & Kanoulas, E. (2020). When to stop reviewing in technology-assisted reviews: Sampling from an adaptive distribution to estimate residual relevant documents. *ACM Transactions on Information Systems (TOIS), 38*(4), 1–36.

Li, H., Scells, H., & Zuccon, G. (2020). Systematic review automation tools for end-to-end query formulation. *Proceedings of the 43rd international acm sigir conference on research and development in information retrieval* (pp. 2141–2144).

Li, H., Wang, S., Zhuang, S., Mourad, A., Ma, X., Lin, J., & Zuccon, G. (2022). To interpolate or not to interpolate: Prf, dense and sparse retrievers. *arXiv preprint arXiv: 2205.00235*.

Lin, J., Nogueira, R., & Yates, A. (2021). Pretrained transformers for text ranking: Bert and beyond. *Synthesis Lectures on Human Language Technologies, 14*(4), 1–325.

Liu, Y.-H. (2009). *The impact of MeSH (Medical Subject Headings) terms on information seeking effectiveness*. Rutgers University-Graduate School-New Brunswick. Ph.D. thesis.

Liu, Y.-H., & Wacholder, N. (2017). Evaluating the impact of MeSH (medical subject headings) terms on different types of searchers. *Information Processing & Management, 53*(4), 851–870.

Lu, Z., Kim, W., & Wilbur, W. J. (2009). Evaluation of query expansion using MeSH in pubmed. *Information Retrieval, 12*(1), 69–80.

Ma, X., Sun, K., Pradeep, R., & Lin, J. (2021). A replication study of dense passage retriever. *arXiv preprint arXiv:2104.05740*.

Moen, S., & Ananiadou, T. S. S. (2013). Distributional semantics resources for biomedical text processing. *Proceedings of LBM*, 39–44.

Nahin, A. (2003). Change to pubmed's automatic term mapping affects phrase searching. *NLM Tech Bull, 331*.

Qin, X., Liu, J., Wang, Y., Liu, Y., Deng, K., Ma, Y., … Sun, X. (2021). Natural language processing was effective in assisting rapid title and abstract screening when updating systematic reviews. *Journal of Clinical Epidemiology, 133*, 121–129. https://doi.org/10.1016/j.jclinepi.2021.01.010

Raffel, C., Shazeer, N., Roberts, A., Lee, K., Narang, S., Matena, M., Zhou, Y., Li, W., & Liu, P. J. (2019). Exploring the limits of transfer learning with a unified text-to-text transformer. *arXiv preprint arXiv:1910.10683*.

Richter, R. R., & Austin, T. M. (2012). Using MeSH (medical subject headings) to enhance pubmed search strategies for evidence-based practice in physical therapy. *Physical Therapy, 92*(1), 124–132.

Sayers, E. (2010). A general introduction to the e-utilities. *Entrez Programming Utilities Help [Internet]. Bethesda: National Center for Biotechnology Information*.

Scells, H., Locke, D., & Zuccon, G. (2018). An information retrieval experiment framework for domain specific applications. *The 41st international acm sigir conference on research & development in information retrieval*.

Scells, H., & Zuccon, G. (2018). searchrefiner: A query visualisation and understanding tool for systematic reviews. *Proceedings of the 27th acm international conference on information and knowledge management* (pp. 1939–1942). ACM.

Scells, H., Zuccon, G., & Koopman, B. (2019). Automatic boolean query refinement for systematic review literature search, . In *WebConf '19The web conference* (pp. 1646–1656).

Scells, H., Zuccon, G., & Koopman, B. (2020a). A computational approach for objectively derived systematic review search strategies. *Proceedings of the 42nd european conference on information retrieval*.

Scells, H., Zuccon, G., & Koopman, B. (2021). A comparison of automatic boolean query formulation for systematic reviews. *Information Retrieval Journal, 24*(1), 3–28.

Scells, H., Zuccon, G., Koopman, B., & Clark, J. (2020b). Automatic boolean query formulation for systematic review literature search. *Proceedings of the web conference 2020* (pp. 1071–1081).

Scells, H., Zuccon, G., Koopman, B., Deacon, A., Azzopardi, L., & Geva, S. (2017). A test collection for evaluating retrieval of studies for inclusion in systematic reviews. *Proceedings of the 40th international acm sigir conference on research and development in information retrieval* (pp. 1237–1240).

Schardt, C., Adams, M. B., Owens, T., Keitz, S., & Fontelo, P. (2007). Utilization of the pico framework to improve searching pubmed for clinical questions. *BMC Medical Informatics and Decision Making, 7*(1), 16.

Schulz, S., Honeck, M., & Hahn, U. (2001). Indexing medical www documents by morphemes. *Studies in Health Technology and Informatics*, (1), 266–270.

Smith, A. M. (2004). An examination of pubmed's ability to disambiguate subject queries and journal title queries. *Journal of the Medical Library Association, 92*(1), 97.

Sneyd, A., & Stevenson, M. (2021). Stopping criteria for technology assisted reviews based on counting processes. *Proceedings of the 44th international acm sigir conference on research and development in information retrieval* (pp. 2293–2297).

Tenopir, C. (1985). Full text database retrieval performance. *Online Review*.

Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł., & Polosukhin, I. (2017). Attention is all you need. *Advances in Neural Information Processing Systems, 30*.

Wacholder, N., Ravin, Y., & Choi, M. (1997). Disambiguation of proper names in text. *Proceedings of the fifth conference on applied natural language processing* (pp. 202–208). Association for Computational Linguistics.

Wang, S., Li, H., Scells, H., Locke, D., & Zuccon, G. (2021a). MeSH term suggestion for systematic review literature search, . In *ADCS '21Proceedings of the 25th australasian document computing symposium*. New York, NY, USA: Association for Computing Machinery. https://doi.org/10.1145/3503516.3503530

Wang, S., Scells, H., Clark, J., Koopman, B., & Zuccon, G. (2022a). From little things big things grow: A collection with seed studies for medical systematic review literature search. *arXiv preprint arXiv:2204.03096*.

Wang, S., Scells, H., Mourad, A., & Zuccon, G. (2022b). Seed-driven document ranking for systematic reviews: A reproducibility study. *European conference on information retrieval* (pp. 686–700). Springer.

Wang, S., Zhuang, S., & Zuccon, G. (2021b). Bert-based dense retrievers require interpolation with bm25 for effective passage retrieval. *Proceedings of the 2021 acm sigir international conference on theory of information retrieval* (pp. 317–324).

Yang, E., MacAvaney, S., Lewis, D. D., & Frieder, O. (2022). Goldilocks: Just-right tuning of bert for technology-assisted review. In M. Hagen, S. Verberne, C. Macdonald, C. Seifert, K. Balog, K. Nørvåg, & V. Setty (Eds.), *Advances in information retrieval* (pp. 502–517). Cham: Springer International Publishing.

Zieman, Y. L., & Bleich, H. L. (1997). Conceptual mapping of user's queries to medical subject headings. *Proceedings of the amia annual fall symposium* (p. 519). American Medical Informatics Association.